

# Gradient Boosted Regression Trees




Material: <https://github.com/pprett/pydata-gbrt-tutorial>


Peter Prettenhofer (@pprett)  
*DataRobot*

Gilles Louppe (@glouppe)  
*Université de Liège, Belgium*



# Motivation



**GECom2012**  
**Wind Forecasting**



# Motivation



# Outline

- 1 Basics
- 2 Gradient Boosting
- 3 Gradient Boosting in scikit-learn
- 4 Use Case: California Housing

# About us

## Peter

- @pprett
- Python & ML  $\sim$  6 years
- sklearn dev since 2010

## Gilles

- @glouppe
- PhD student (Liège, Belgium)
- sklearn dev since 2011  
*Chief tree hugger*



# Outline

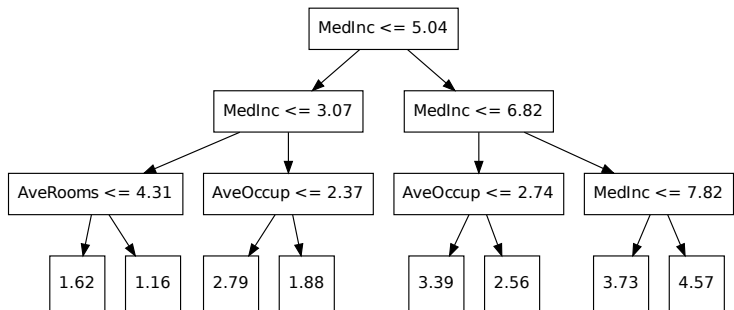
- 1 Basics
- 2 Gradient Boosting
- 3 Gradient Boosting in scikit-learn
- 4 Use Case: California Housing

# Machine Learning 101

- Data comes as...
  - A set of examples  $\{(\mathbf{x}_i, y_i) | 0 \leq i < n\_samples\}$ , with
  - Feature vector  $\mathbf{x} \in \mathbb{R}^{n\_features}$ , and
  - Response  $y \in \mathbb{R}$  (regression) or  $y \in \{-1, 1\}$  (classification)
- Goal is to...
  - Find a function  $\hat{y} = f(\mathbf{x})$
  - Such that error  $L(y, \hat{y})$  on new (unseen)  $\mathbf{x}$  is minimal

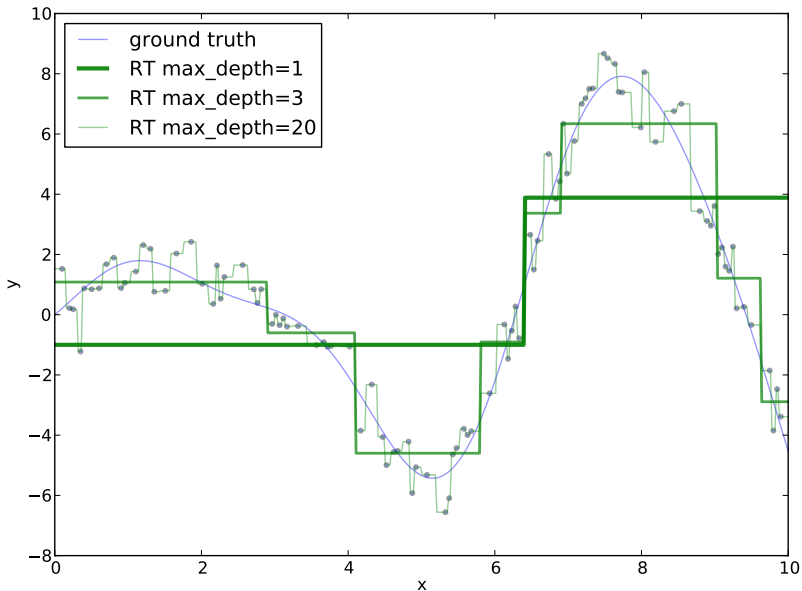


# Classification and Regression Trees [Breiman et al, 1984]

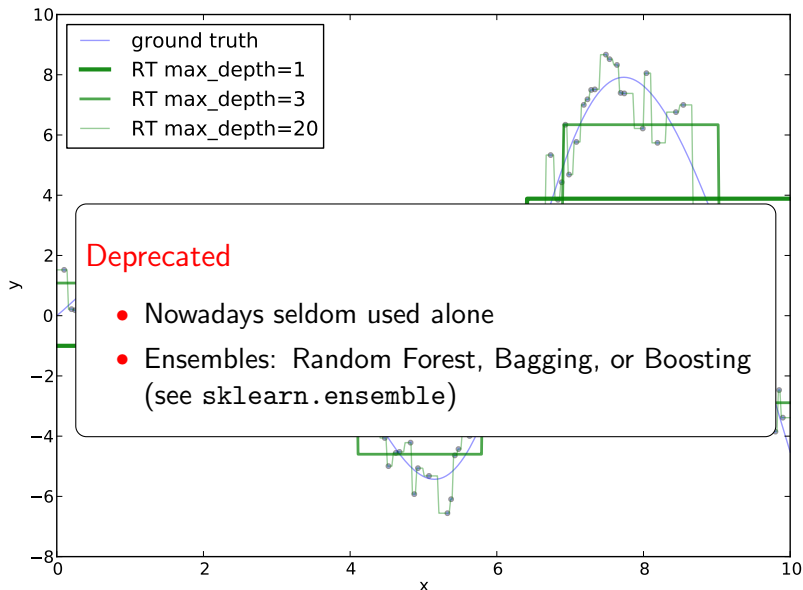




# Function approximation with Regression Trees



# Function approximation with Regression Trees



# Outline

- 1 Basics
- 2 Gradient Boosting**
- 3 Gradient Boosting in scikit-learn
- 4 Use Case: California Housing


# Gradient Boosted Regression Trees

## Advantages

- Heterogeneous data (features measured on different scale)
- Supports different loss functions (e.g. huber)
- Automatically detects (non-linear) feature interactions

## Disadvantages


- Requires careful tuning
- Slow to train (but fast to predict)
- Cannot extrapolate



# Boosting

## AdaBoost [Y. Freund & R. Schapire, 1995]


- Ensemble: each member is an expert on the errors of its predecessor
- Iteratively re-weights training examples based on errors



# Boosting

## Huge success

- Viola-Jones Face Detector (2001)



- Freund & Schapire won the Gödel prize 2003

# Gradient Boosting [J. Friedman, 1999]

Statistical view on boosting


- $\Rightarrow$  Generalization of boosting to arbitrary loss functions

# Gradient Boosting [J. Friedman, 1999]

## Statistical view on boosting

- $\Rightarrow$  Generalization of boosting to arbitrary loss functions

## Residual fitting





# Functional Gradient Descent

## Least Squares Regression

- Squared loss:  $L(y_i, f(\mathbf{x}_i)) = (y_i - f(\mathbf{x}_i))^2$
- The residual  $\sim$  the (negative) gradient  $\frac{\partial L(y_i, f(\mathbf{x}_i))}{\partial f(\mathbf{x}_i)}$



# Functional Gradient Descent

## Least Squares Regression

- Squared loss:  $L(y_i, f(\mathbf{x}_i)) = (y_i - f(\mathbf{x}_i))^2$
- The residual  $\sim$  the (negative) gradient  $\frac{\partial L(y_i, f(\mathbf{x}_i))}{\partial f(\mathbf{x}_i)}$

## Steepest Descent

- Regression trees approximate the (negative) gradient
- Each tree is a successive gradient descent step



# Outline

- 1 Basics
- 2 Gradient Boosting
- 3 Gradient Boosting in scikit-learn**
- 4 Use Case: California Housing

## Notebook

<https://github.com/pprett/pydata-gbrt-tutorial>

# Outline

- 1 Basics
- 2 Gradient Boosting
- 3 Gradient Boosting in scikit-learn
- 4 Use Case: California Housing